

## CS33 Lecture 1: Bits and Bytes

- Pre-processor  $\rightarrow$  compiler  $\rightarrow$  assembler  $\rightarrow$  linker
  - $\hookrightarrow$  text to binary thru assembler
  - $\hookrightarrow$  assembler is machine specific
    - $\hookrightarrow$  on a binary level, program is run for a particular machine architecture
- Hex-denoted by 0x preceding hex number
- Send help
- Left shift  $\rightarrow$  shift bit vector left, fill w/ zeroes
- Right Shift
  - $\hookrightarrow$  Arithmetic  $\rightarrow$  shift bit vector right, replicate most significant digit
  - $\hookrightarrow$  Logical  $\rightarrow$  shift bit vector right, fill w/ zeroes

## CS33 Lecture 2: Integers

### • Encoding integers

↳ Unsigned:  $B2U(x) = \sum_{i=0}^{n-1} x_i \cdot 2^i$

↳ Two's Complement:  $B2T(x) = -x_{n-1} \cdot 2^{n-1} + \sum_{i=0}^{n-2} x_i \cdot 2^i$

↳ most significant bit indicates sign

↳ 0000 0001 0010 0011 . . . 0111 1000 . . . 1111

U: 0 1 2 3 . . . 7 8 . . . 15

T: 0 1 2 3 . . . 7 -8 . . . -1

### ↳ Maxes/Mins

↳  $|T_{\min}| = T_{\max} + 1 \rightarrow$  asymmetric range

↳  $U_{\max} = 2^{\lfloor \log_2 n \rfloor}$

↳ `#include <limits.h>` for C Programming

↳ constants: `ULONG_MAX, LONG_MAX, LONG_MIN`

↳ Word Size - size used for pointers (ie 64-bit machine encodes pointers as 64-bit)

### ↳ Casting

↳ signed values implicitly cast to unsigned

↳  $-1 > 0U \rightarrow -1$  unsigned is  $U_{\max}$

↳  $2147483647U < -2147483647 - 1 \rightarrow 011..11 < 11..110$

↳ `(unsigned)-1 > -2`

↳  $2147483647 < 2147483648U \rightarrow T_{\max} < T_{\max} + 1$

↳  $2147483647 > (int)2147483648U$

## CS33 Lecture 3: Machine-Level Programming: Basics

- CPU puts addresses into memory
  - ↳ Memory returns data/instructions
- Register is part of CPU
  - ↳ faster than memory (SRAM vs. DRAM)
  - ↳ compiler-visible like memory
- PC - program counter - tells us what instruction is being executed
- Registers and memory together hold the state of the program
- Memory is byte-addressable
- x86-64 has 16 Integer Registers
  - ↳ %rax has 64 bits
    - ↳ contains %eax → lower 32-bits, can reference individually
    - ↳ can reference lower, but not upper
  - ↳ %rsp points to the current location in memory where the stack is
  - ↳ Memory is contiguous in each register file
- Moving Data
  - ↳ movq → copies from source → destination
  - ↳ operand can be 1 of 3 types:
    - ↳ immediate - constant integer data
      - ↳ Ex) \$0x400, 1-533
    - ↳ register - one of 16 integer registers
      - ↳ Ex) %rax, %r13 → location, not value
    - ↳ memory - 8 bytes of memory at address given by register
      - ↳ Ex) (%rax) → pointer for movq instruction (dereference)
  - ↳ imm → reg: movq \$0x4, %rax → temp = 0x4
  - ↳ imm → mem: movq \$-147, (%rax) → \*p = -147
  - ↳ reg → reg: movq %rax, %rdx → temp2 = temp1
  - ↳ reg → mem: movq %rax, (%rdx) → \*p = temp
  - ↳ mem → reg: movq (%rax), %rdx → temp = \*p
    - ↳ mem denoted with parens

### • Complete Memory Addressing Modes

- ↳ D(Rb, Ri, S) → Mem[Reg(Rb) + S \* Reg(Ri) + D]
  - ↳ Ex) 8(%ax, %dx, 2) → accesses rax + 2 \* rdx + D location

## • Address Computation Instruction

- ↳ leaq → doesn't dereference, doesn't go into memory
- ↳ setting up an address (finding a pointer)
- ↳ leaq (rax), rdx vs. movq (rax), rdx
- ↳ leaq (rax), rdx → movq rax, rdx
- ↳ leaq (rax, rbx), rdx → add rax and rbx and write it to rdx
- ↳ useful for 2-address code (an address is an input and output)

## CS33 Lecture 4: Machine-Level Programming - Control

- Condition codes: 1-bit registers

- CF - carry flag (unsigned)

- ZF - zero flag

- SF - sign flag (signed)

- OF - overflow flag (signed)

- Ex) using addq:  $t = a + b$

- CF set if MSB carried out

- ZF set if  $t = 0$

- SF set if  $t < 0$  (set if signed or unsigned)

- OF set if 2s complement overflow " → set w/ CF

- not set by leaq \*

- cmpq instruction ( $\text{Src2}, \text{Src1}$ ) → ( $\text{Src1} - \text{Src2}$ )

- performs subtraction without setting destination

- saves register space when performing comparisons

- CF: carry out, ZF:  $a = b$ , SF:  $(a - b) < 0$ , OF: 2s complement overflow

- Jumping - jump to diff. parts of code depending on condition codes

- can be set explicitly (cmpq) or implicitly (arithmetic)

- modifies RIP → points to current instruction (instruction pointer)

- jmp → unconditional jump. UBR

- rest of instructions are CBR → conditional

- Ex) je/jne → jumps if equal/zero or notequal/nonzero

- Ex) ja/jb → unsigned

- check if  $< 0$  or  $> 10 \rightarrow$  can use ja → all neg. values interpreted as  $> 10$

- linker replaces labels with memory addresses for jumping

- Conditional move: using more registers, allows us to avoid control flow issues

- Ex) if( $x < y$ ) → return  $x - y$ , else return  $y - x$

- writes both results into registers before evaluating condition

- decides which to return based on condition

- compiler stores alternatives that will be performed to be returned/selected later

- cmov → conditional move instruction

- Bad when expensive (hard evaluations), risky (dereferencing), side effects present (value modification)

• All loops can generally be modified to the same structure

↳ Do-while

↳ .L2:

movq rdi, edx → rdi is x

andl \$1, edx → edx is either 0 or 1

addq rdx, rax

shrq rdx, rax → sets ZF when x=0, x>>=1

jne .L2

rep; ret

↳ Switch

↳ case 100:

i++; j++;

break;

case 102:

i+=2; j+=2;

break;

case 107: ← fall through

i++;

case 104:

j+=2;

break;

case 106:

i+=2;

j++;

, break;

default:

i=0; j=0;

↳ jump table (values close together)

↳ one jump instruction sends the RIP to find the correct memory address

↳ cases merge back to the break

• Direct jump → jbe → to a location

• Indirect jump → jmpq \*(ADDR) → used in the jump table

## CS33 Lecture 5: Machine-Level Programming: Procedures

- The stack grows down in addresses ( $FFF \rightarrow 000$ )
  - ↳  $rsp$  points to the top of the stack (lowest address)
  - ↳ all programs have their own stack/heap  $\rightarrow$  virtual memory
- Ex) 400540 <multstore>:
  - :
  - 400544: callq 400550 <mult2>
  - 400549: mov %rax, (%rbx)
  - :
  - 400550 <mult2>:
  - 400550: mov %rdi, %rax
  - :
  - 400557: retq

- ↳ rip increments with the instruction addresses
- ↳ naturally goes to the next instruction
  - ↳ changes w/ jumps, callq, etc.
- ↳ callq leaves the original caller's memory address on the stack and changes rip to the new instruction address
- ↳ retq  $\rightarrow$  pop + jump  $\rightarrow$  consumes data in the stack, changes rip back to caller's address
  - ↳ uses the  $rsp$  to locate the caller's address

### • Procedure Data Flow

- ↳ 1st 6 arguments stored in registers

- ↳  $rdi, rsi, rdx, rcx, r8, r9$

- ↳ 7+ arguments placed on stack

### • x86-64 / Linux Stack Frame

- ↳ Frame-data of some function call

- ↳ Arguments (if 7+), return address  $\rightarrow$  callee frame

- ↳ callee frame: Registers + vars, argument build (if caller of another function)

### • Register Saving Conventions

- ↳ Caller saved  $\rightarrow$  temp. values saved in frame before call

- ↳ Callee saved  $\rightarrow$  temp. values saved in frame after call, before using

- ↳ restored by callee before returning to caller

- x86-64 Linux Register Usage
  - ↳ rax → return value
  - ↳ caller saved, can be modified by procedure
  - ↳ rdi, rsi, rdx, rcx, r8, r9 → arguments
  - ↳ caller saved, can be modified by procedure
  - ↳ r10, r11
  - ↳ caller saved
  - ↳ rbx, r12, r13, r14 → callee saved
  - ↳ rbp → callee saved
    - ↳ may be used as a frame-pointer
  - ↳ rsp → special form of callee save, stack pointer

## CS33 Lecture 6: Machine-Level Programming: Data

### • Multi-dimensional Arrays

↳ Declaration:  $T A[R][C]$

↳ 2D array of type  $T$

↳ R rows, C columns

↳ Type  $T$  element requires  $K$  bytes

↳ Array Size:  $R * C * K$  bytes

↳ Arrangement: Row-major ordering

### • Multi-level Array

↳ Array with a first array that contains pointers to secondary arrays

↳ More flexible, more complex

↳ Can have varying lengths at each unit

### • Machine-level Translation

↳ Fixed-Size  $\rightarrow C=16, K=4 \rightarrow A + i * (16 * K) + j * K$

↳ a in rdi, i in rsi, j in rdx  $\rightarrow i * (16 * K)$  positions which nested array

↳ salq \$6, %rsi #  $64 * i$   $\rightarrow j * K$  positions inside the nested array

addq %rsi, %rdi #  $a + 64 * i$

movl (%rdi,%rdx,4), %eax #  $M[a + 64 * i + 4 * j]$

ret

↳ Dynamic  $\rightarrow C=n, K=4 \rightarrow A + i * (C * K) + j * K$

↳ must use multiplication rather than shifts

↳ n in rdi, a in rsi, i in rdx, j in rcx

↳ imulq %rsi, %rdi #  $n * i$

leaq (%rsi,%rdi,4), %rax #  $a + 4 * n * i$

movl (%rax,%rcx,4), %rax #  $a + 4 * n * i + 4 * j$

ret

### • Structs

↳ must satisfy alignment requirement  $K = \text{largest alignment of any element}$

↳ prevents overlapping of blocks/pages/etc

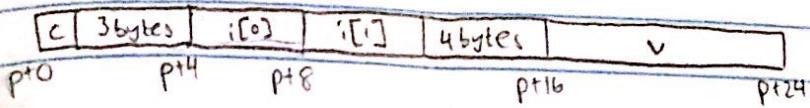
↳ struct S1 {

char c;

int i[2];

double v;

}



## CS33 Lecture 8: Floating-Points

### - Tiny FP example (8 bits)

↳ the sign is in the MSB

↳ next 4 bits are the exponent  $\rightarrow$  bias of 7 ( $2^4 - 1$ )

↳ the last 3 bits are the frac

↳ normalized:  $(-1)^S \cdot 1.\text{frac} \times 2^E$ ,  $E = \text{exp-bias}$

↳ exp! = 000...0, exp! = 111...1

↳ denormalized  $\rightarrow$  exp = 000...0

↳  $E = 1\text{-Bias}$ ,  $M = 0.\text{frac}$

↳ infinity: exp is 111...1, frac is zero

↳ NaN: exp is 111...1, frac is non-zero

↳ instead of overflow, floating-point arithmetic saturates

•  $x \neq \pm (\text{int})(\text{float})$   $x \rightarrow$  float doesn't have enough precision

•  $x = (\text{int})(\text{double})$   $x \checkmark$

•  $f = \pm (\text{float})(\text{double})$   $f \checkmark$

•  $\delta = \pm (\text{double})(\text{float})$   $\delta$

•  $f = -(-f) \checkmark \rightarrow$  float doesn't have the asymmetry of int

•  $2/3 \neq 2/3.0$

•  $\delta < 0.0 \rightarrow ((\delta * 2) < 0.0) \checkmark \rightarrow$  no overflow concerns

•  $\delta > f \rightarrow -f > -\delta \checkmark$

•  $\delta * \delta \geq 0.0 \rightarrow$  no overflow concerns

•  $(\delta + f) - \delta \neq f \rightarrow$  potential for not enough precision, f rounded away

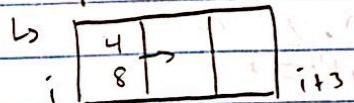
## CS33 Lecture 9: Optimization

Feb 6 - 10, 2024

- The compiler has to be conservative
  - ↳ Branching can block optimization → not sure if code is actually executed
  - ↳ Storing to memory can block optimization → memory alteration
  - ↳ Functions can block memory
    - ↳ Inlining can help, but you lose access to interposition (updating with DLLs without recompiling)
    - ↳ Interposition leaves a gap in the code assembly
    - ↳ The compiler cannot assume the function doesn't modify memory
- Simple optimizations: code motion, strength reduction, subexpressions
- Memory conflicts: aliasing
  - ↳ No guarantee that two references to memory don't overlap → cannot optimize
  - ↳ Introducing local variables → tells compiler to not check for aliasing

### -Architecture:

- ↳ Processor can execute instructions in diff. order than compiler, must affect architectural units in correct order
- ↳ Improves performance, reordered at end by retirement unit
- ↳ Hardware does a lot of operation
- ↳ Multiple functional units → instruction-level parallelism



↳

↳ Computation divided into stages, independent computations can be passed from stage to stage

↳ Latency bound vs. throughput bound → time vs. # of tasks

↳ A straight-line chain of computation is not optimal

### • SIMD - Single Instruction Multiple Data

↳ Vector elements → vector registers → hold multiple values

• Branch Prediction → attempt to predict on branching, penalty when wrong prediction

↳ must flush pipeline if wrong, high cycle cost

↳ predictable branches have lower overhead

## CS33 Lecture 10: Memory Hierarchy

- CPU sends an address out, receives data/instructions from memory
  - ↳ Transistors scale well, memory wall means memory access doesn't
  - ↳ CPU can only be as fast as its ability to access memory
  - ↳ Bound by memory

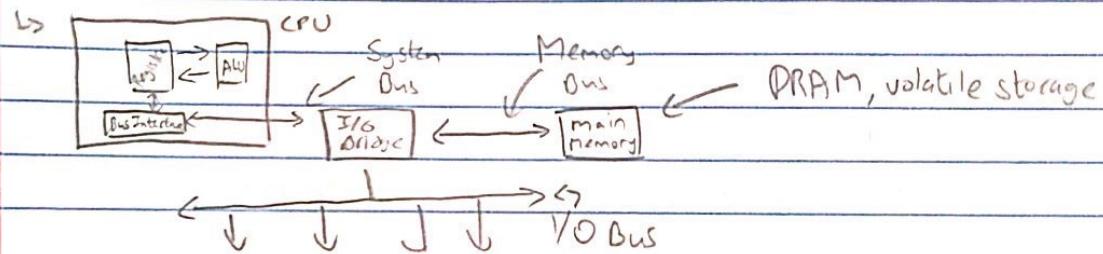
- SRAM - Static Random Access Memory → registers, caches, etc

- ↳ fast to access, hard to have high capacity

- DRAM - Dynamic Random Access Memory → main memory

- ↳ slower to access higher capacity

- I/O Bus



- ↳ non-volatile storage: disks, USB drive, etc.

- ↳ CPU request to disk, disk stored in main memory (DMA), CPU notified by interrupt

- Costs: area, material, power, monetary, etc.

- Locality: Programs tend to use data and instructions with addresses near those they have used recently

- ↳ Temporal: recently referenced items are likely to be referenced again in the near future (loop w/ iteration)

- ↳ Spatial: items with nearby addresses tend to be referenced close together (arrays and structs)

- General Cache Concepts

- ↳ Cache → smaller, faster, more expensive memory caches a subset of blocks transferred from Main memory

- ↳ brings reused data in to lower latency of common accesses

- ↳ blocks may contain larger granularities than necessary

- ↳ why padding is necessary, takes advantage of spatial locality

- ↳ decisions based on cache-policy combined w/ memory stream (eviction)

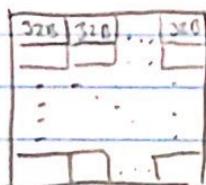
- Prefetching - attempt to access data ahead of time

- ↳ further ahead of time = lower latency, balance w/ pollution of storage

## CS33 Lecture 11: Memory and Parallelism

### • Matrix Multiplication

↳



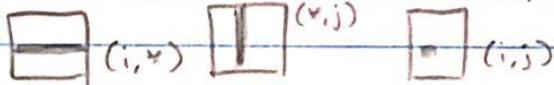
↳ B

81 81 81  
81 81 81  
81 81 81

→ machine manipulated, not compiler

→ about reuse for locality

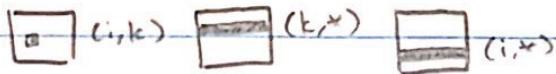
↳ 1st method → k in inner-most loop



↳ row order with i → 0.25 misses/iteration

↳ column order with j → 1 miss/iteration

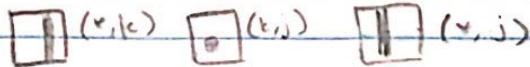
↳ 2nd method → j in inner-most loop



↳ row-order w/ B and C → 0.25 misses/iteration

↳ more operations, better locality

↳ 3rd method → i in inner-most loop



↳ worst, A and C miss every iteration

### • Blocked Matrix Multiplication

↳ Constrain matrix to work in smaller matrices

$$C = \boxed{A} * \boxed{B} + \boxed{C}$$

↳ multiple rows/columns per block

↳ constraining amt of data accessed → can control how much data the cache deals w/

↳ can be worse due to fine complexity / branching

### • Cache Miss Analysis

↳ Assume → block = 8 doubles,  $C << n$ , 3 blocks in cache ( $3B^2 < C$ )

↳  $B^2/8$  misses for each block

$$\Rightarrow \frac{2n}{B} * \frac{B^2}{8} = \frac{nB}{4}$$

### • MIMD

↳ Multi-core → simultaneous execution when data sharing is unnecessary

↳ Cooperative multithreading → separate execution, but sharing of some data

## CS33 Lecture 11: Parallelism Part 1

- $T_0 \quad T_1 \quad T_2 \quad T_3 \rightarrow$  ideally Parallel Latency = Latency/4
- ↳ Sometimes less performance
  - ↳ overhead  $\rightarrow$  communication between cores / OS interaction
  - ↳ ILP  $\rightarrow$  may not be much parallelism in task
  - ↳ load balancing  $\rightarrow$  work can't be distributed evenly
- ↳ Sometimes more performance:
  - ↳ Memory bound applications can get much better memory locality by distributing across more caches
  - ↳ Balance computation v. communication
- Work sharing
- ↳ for:
  - ↳ `#pragma omp parallel`
  - ⋮
  - ↳ `#pragma omp for`
  - ↳ `for(i=0, i<10000; i++)`  $\rightarrow$  Static
  - $\begin{array}{c} \checkmark \quad \triangleright \\ T_0 \quad T_1 \\ 5000 \quad 5000 \end{array}$
  - ↳ static  $\rightarrow$  tasks split, no further communication / load balancing
  - ↳ dynamic  $\rightarrow$  requests more work when done with work
    - ↳ much more overhead due to communication
  - ↳ guided  $\rightarrow$  start off with large chunks of work distributed and size decreases logarithmically
  - ↳ cannot have breaks  $\rightarrow$  openmp must know at runtime how many iterations the loop will take

### ↳ single:

- ↳ communicates that a piece of code is only to be done by 1 thread
- ↳ critical sections  $\rightarrow$  only one thread can access this at a time  $\rightarrow$  sequential
  - ↳ a lot of overhead in creating this region
- ↳ reduction:
  - ↳ specifies associative operator and the variable
  - ↳ private variable created in each thread, after running, the associative operator is used on the variables to combine
- Deadlock  $\rightarrow$  mutually exclusive resource access, threads hold resources until they have what they want, resources cannot be taken away, cyclical

## CS 33 Lecture 12: Parallelism Part 2

- 2 types of parallelism

- ↳ Data → implicit parallelism inside data itself

- ↳ Task → specialization of task across cores

- ↳ executes different things on each task

- OpenMP Tasks → independent units of work

- ↳ Composed of code, data, and internal control variables

- ↳ Threads perform the work of each task

- ↳ Runtime system decides execution → may be deferred, or executed immediately

- ↳ guaranteed completion at barriers

- ↳ barrier → thread barrier

- ↳ taskwait → task barrier

- Scope - variables are default private in a task

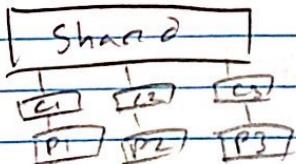
- ↳ shared → makes variables shared between threads

- ↳ firstprivate → makes sure variable is private / has the correct value

- OpenMP memory model

- ↳ supports a shared memory model → threads share an address space

- ↳



- ↳ defined by

- ↳ coherence → behavior when single address is accessed by multiple threads

- ↳ consistency → orderings of reads, writes, syncs w/ various addresses and multiple threads

- ↳ weak consistency → can't reorder sync ops w/ read or write ops

## CS33 Lecture 14 : Exceptional Control Flow and Linking

- Synchronous Exceptions

- ↳ Traps

- ↳ intentional → part of the actual code

- ↳ system calls, breakpoint traps, special instructions

- ↳ returns control to "next" instruction

- ↳ Fault

- ↳ unintentional but possibly recoverable

- ↳ page faults, protection faults, floating point exceptions

- ↳ either re-executes faulting instruction or aborts

- ↳ Aborts

- ↳ unintentional and unrecoverable

- ↳ illegal instruction, parity error, machine check

- ↳ aborts current program

- Page faults

- ↳ user writes to memory location, which is currently on disk

- ↳ OS copies page from disk to memory

- Invalid Memory Reference - fault

- ↳ starts w/ page fault, OS detects invalid address, signals the process

- ↳ sends SIGSEGV, program exits w/ segmentation-fault

- System Calls - trap

- ↳ each x86-64 system call has a unique ID #

- ↳ OS needs to intervene (read, write, open, etc.)

- ↳ sends exception w/ ID # in eax, performs the specified instruction, returns with status of operation

- Compilation:

- ↳ CPP → preprocessor → #defines, #pragmas → injects source code

- ↳ compiler → parses → forms an intermediate representation of source code across diff. languages → optimizes code

- ↳ assembler → change into binary

- ↳ multiple object files, may have unresolved symbols, incomplete view of memory layout

- ↳ linker → resolves above issues

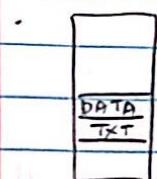
- Creating a library

- ↳ Put all functions in a single source file → space/time inefficient

- ↳ Put each function in a separate source file → burdensome on programmer

- ↳ related object files into a single archive file, linker searches thru .a for symbols

## CS33 Lecture 15: Virtual Memory



32-bit

$2^{32}$  B

↳ Prog  $\xrightarrow{\text{Addr}}$  Mem  
Data

↳ Model cannot handle simultaneous, independently compiled programs

↳ wastes memory, must account for all possible combinations

↳ VM: Prog  $\xrightarrow{\text{VA}}$  MMU  $\xrightarrow{\text{PA}}$  Mem.

↳ Applications have own view of memory, translated into single set of main mem.

↳ Excess memory not in DRAM, in disk  $\rightarrow$  exploit locality

• Page Table - takes VA, translates into PA  $\rightarrow$  dynamic by necessity

↳ resides in physical memory - may not fit in memory as a whole

↳ VA  $\xrightarrow{\text{PT}}$  PA

↳ Communication between VM and MM done in pages

↳ Relatively high access cost  $\rightarrow$  common swaps = thrashing

↳ A page table register per process

↳ can protect from duplicate addresses across processes or allow sharing

↳ Page table start denoted by page table base register (PTBR)

↳ Each entry has a valid bit that tells if the entry is in physical memory

↳ accessing an entry with an invalid bit results in a page fault  $\rightarrow$

↳ evicts old page from memory and replaces w/ requested page

↳ behavior is non-deterministic

↳ Dirty page  $\rightarrow$  page in memory doesn't match page on disk  $\rightarrow$  consistency

↳ Breakdown of address

↳ Page offset  $\rightarrow$  index within page  $\rightarrow$  unchanged between VA and PA

↳ Page number  $\rightarrow$  specifies which page in VM is translated into Mem

• TLB - Translation Lookaside Buffer

↳ Hardware structure acts as a cache for the page table

↳ Takes you to physical address directly

- Multi-Level Page Tables

- ↳ Level 1 → each PTE points to a pagetable
  - ↳ alias for only part of page table in DRAM
  - ↳ more loads → less severe if good locality

- ↳

VPN1	VPN2	...	VPN $\ell$	VPO
------	------	-----	------------	-----

- Balancing page size, spatial locality, and performance

## CS33 Lecture 1b: MIPS

### CISC v. RISC

- ↳ CISC motivated by capacity of memory limitations
  - ↳ complex instructions → mem. addressing mode
  - ↳ memory access or register address → all in 1 instruction = compact
  - ↳ small # of registers
  - ↳ variable length instructions
  - ↳ difficult to decode / pipeline / implement
- ↳ Compilers becoming more powerful, complexity rising → move to RISC
  - ↳ simpler instructions, fixed length
    - ↳ wastes space, simplifies decoding
  - ↳ boundaries blurred to allow both architectures to coexist

### MIPS

- ↳ Fixed length → 32 bits
- ↳ R type → register specifiers ( $R_s, R_t, R_d$ ) → 5 bit, ALU oriented
  - ↳ 32 registers
  - ↳ dest, src
- ↳ I type - load word and store word → data transfer
  - ↳ load = mem. → reg., store = reg. → mem.
  - ↳ LW \$t0, 4(\$s0) →  $R[t_0] \leftarrow M[R[s_0] + 4]$  (src, dest)
    - ↳ Base(register) + displacement (immediate)
- ↳ J type - jump(j) and jump and link(jnl)
  - ↳ j changes the PC
  - ↳ jnl is like a callq → \$RA set to PC+4, PC to new PC
    - ↳ nested calls → RA saved on the stack

## CS33 Lecture 17: MIPS

- Pseudo instructions

↳ translated into other machine instructions

↳ move → move \$t,\$s → addiu \$t,\$s,0

↳ clear → clear \$t → addu \$t,\$zero,\$zero

↳ load 16-bit i → li \$t,C → addiu \$t,\$zero,C-10

↳ load 32-bit i → li \$t,C → lui \$t,C-hi,ori \$t,\$t,C-lo

↳ load label address → la \$t,A → lui \$t,A-hi,ori \$t,\$t,A-lo

↳ instructions have 6 bits for op-code, 5 for each register → only 16 (2)  
hold immediate → must separate C-hi and C-lo

- System calls → traps are recoverable and intentional

↳ application trying to perform I/O that requires communication w/ OS →  
syscall

↳ passes a code as a parameter to identify the specific syscall → \$V0

↳ switch statement - case

· \$a0 and \$a1 hold parameters

LH SC sub(\$t) \$a0

addu \$t,\$a0,\$a1

addiu \$t,\$a0,\$a1